

Practical Machine Learning Course Project

Luca Caramellino

Synopsis

Use data from accelerometers wore by volunteer participants executing barbell to identify the quality of the exercise. Some volunteer executed the exercise correctly while other purposely committed common mistakes. The predictive software must be able to identify the correct exercises with reasonable accuracy.

Load libraries

```
library(AppliedPredictiveModeling)
library(caret)
library(rattle)
library(rpart.plot)
library(randomForest)
library(knitr)
library(e1071)
```

Load Data

```
training <- read.csv("pml-training.csv", na.strings=c("NA", ""), header=TRUE)
colnames_train <- colnames(training)
testing <- read.csv("pml-testing.csv", na.strings=c("NA", ""), header=TRUE)
colnames_test <- colnames(testing)
```

Filter Data

Data sets are filtered to remove NA values and near zero variables.

```
nonNAs <- function(x) {
  as.vector(apply(x, 2, function(x) length(which(!is.na(x)))))
}

colcnts <- nonNAs(training)
drops <- c()
for (cnt in 1:length(colcnts)) {
  if (colcnts[cnt] < nrow(training)) {
    drops <- c(drops, colnames_train[cnt])
  }
}

training <- training[,!(names(training) %in% drops)]
training <- training[,8:length(colnames(training))]

testing <- testing[,!(names(testing) %in% drops)]
testing <- testing[,8:length(colnames(testing))]
```

The training set has **19622** samples and **52** potential predictors after filtering. The testing set result instead with **20** samples and **52** predictors.

Check for Covariates

```
nsv <- nearZeroVar(training, saveMetrics=TRUE)
nsv
```

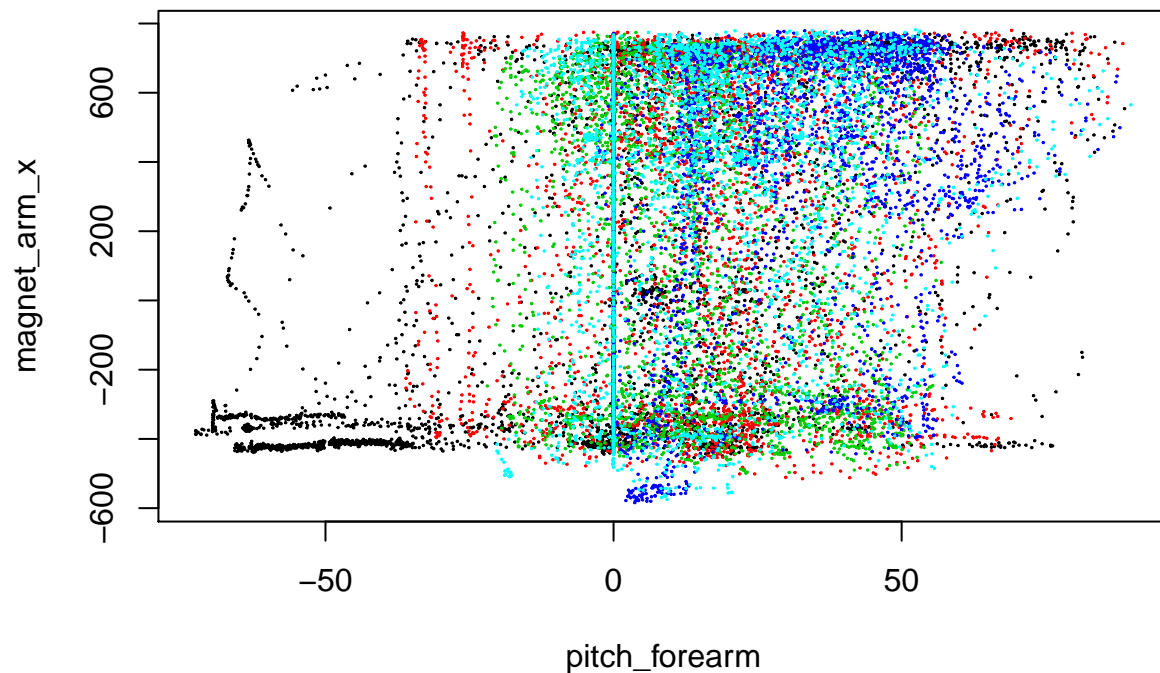
##	freqRatio	percentUnique	zeroVar	nzv
## roll_belt	1.101904	6.7781062	FALSE	FALSE
## pitch_belt	1.036082	9.3772296	FALSE	FALSE
## yaw_belt	1.058480	9.9734991	FALSE	FALSE
## total_accel_belt	1.063160	0.1477933	FALSE	FALSE
## gyros_belt_x	1.058651	0.7134849	FALSE	FALSE
## gyros_belt_y	1.144000	0.3516461	FALSE	FALSE
## gyros_belt_z	1.066214	0.8612782	FALSE	FALSE
## accel_belt_x	1.055412	0.8357966	FALSE	FALSE
## accel_belt_y	1.113725	0.7287738	FALSE	FALSE
## accel_belt_z	1.078767	1.5237998	FALSE	FALSE
## magnet_belt_x	1.090141	1.6664968	FALSE	FALSE
## magnet_belt_y	1.099688	1.5187035	FALSE	FALSE
## magnet_belt_z	1.006369	2.3290184	FALSE	FALSE
## roll_arm	52.338462	13.5256345	FALSE	FALSE
## pitch_arm	87.256410	15.7323412	FALSE	FALSE
## yaw_arm	33.029126	14.6570176	FALSE	FALSE
## total_accel_arm	1.024526	0.3363572	FALSE	FALSE
## gyros_arm_x	1.015504	3.2769341	FALSE	FALSE
## gyros_arm_y	1.454369	1.9162165	FALSE	FALSE
## gyros_arm_z	1.110687	1.2638875	FALSE	FALSE
## accel_arm_x	1.017341	3.9598410	FALSE	FALSE
## accel_arm_y	1.140187	2.7367241	FALSE	FALSE
## accel_arm_z	1.128000	4.0362858	FALSE	FALSE
## magnet_arm_x	1.000000	6.8239731	FALSE	FALSE
## magnet_arm_y	1.056818	4.4439914	FALSE	FALSE
## magnet_arm_z	1.036364	6.4468454	FALSE	FALSE
## roll_dumbbell	1.022388	84.2065029	FALSE	FALSE
## pitch_dumbbell	2.277372	81.7449801	FALSE	FALSE
## yaw_dumbbell	1.132231	83.4828254	FALSE	FALSE
## total_accel_dumbbell	1.072634	0.2191418	FALSE	FALSE
## gyros_dumbbell_x	1.003268	1.2282132	FALSE	FALSE
## gyros_dumbbell_y	1.264957	1.4167771	FALSE	FALSE
## gyros_dumbbell_z	1.060100	1.0498420	FALSE	FALSE
## accel_dumbbell_x	1.018018	2.1659362	FALSE	FALSE
## accel_dumbbell_y	1.053061	2.3748853	FALSE	FALSE
## accel_dumbbell_z	1.133333	2.0894914	FALSE	FALSE
## magnet_dumbbell_x	1.098266	5.7486495	FALSE	FALSE
## magnet_dumbbell_y	1.197740	4.3012945	FALSE	FALSE
## magnet_dumbbell_z	1.020833	3.4451126	FALSE	FALSE
## roll_forearm	11.589286	11.0895933	FALSE	FALSE
## pitch_forearm	65.983051	14.8557741	FALSE	FALSE
## yaw_forearm	15.322835	10.1467740	FALSE	FALSE
## total_accel_forearm	1.128928	0.3567424	FALSE	FALSE

## gyros_forearm_x	1.059273	1.5187035	FALSE	FALSE
## gyros_forearm_y	1.036554	3.7763735	FALSE	FALSE
## gyros_forearm_z	1.122917	1.5645704	FALSE	FALSE
## accel_forearm_x	1.126437	4.0464784	FALSE	FALSE
## accel_forearm_y	1.059406	5.1116094	FALSE	FALSE
## accel_forearm_z	1.006250	2.9558659	FALSE	FALSE
## magnet_forearm_x	1.012346	7.7667924	FALSE	FALSE
## magnet_forearm_y	1.246914	9.5403119	FALSE	FALSE
## magnet_forearm_z	1.000000	8.5771073	FALSE	FALSE
## classe	1.469581	0.0254816	FALSE	FALSE

No covariates was identified so there is no need to further filtering the data set.

Plot features with highest correlation with classe

```
cor <- abs(sapply(colnames(training[, -ncol(training)]), function(x) cor(as.numeric(training[, x]), as.numeric(training[, "classe"]))))
plot(training[, names(which.max(cor))], training[, names(which.max(cor[-which.max(cor)]))], col = training[, "classe"])
```



There isn't any strong predictors that correlates with `classe` therefore linear regression is not suitable. Random forests algorithm may generate more robust predictions for our data and is therefore selected.

Algorithm (Random Forest)

Creating smaller dataset from the original one

The training dataset is divided into smaller sets both to avoid overfitting and to allow predictive algorithm to run faster.

```
set.seed(3)
ids_small <- createDataPartition(y=training$classe, p=0.25, list=FALSE)
small1 <- training[ids_small,]
remainder <- training[-ids_small,]

set.seed(333)
ids_small <- createDataPartition(y=remainder$classe, p=0.33, list=FALSE)
small2 <- remainder[ids_small,]
remainder <- remainder[-ids_small,]

set.seed(333)
ids_small <- createDataPartition(y=remainder$classe, p=0.5, list=FALSE)
small3 <- remainder[ids_small,]
small4 <- remainder[-ids_small,]

set.seed(333)
inTrain <- createDataPartition(y=small1$classe, p=0.6, list=FALSE)
small_training1 <- small1[inTrain,]
small_testing1 <- small1[-inTrain,]

set.seed(333)
inTrain <- createDataPartition(y=small2$classe, p=0.6, list=FALSE)
small_training2 <- small2[inTrain,]
small_testing2 <- small2[-inTrain,]

set.seed(333)
inTrain <- createDataPartition(y=small3$classe, p=0.6, list=FALSE)
small_training3 <- small3[inTrain,]
small_testing3 <- small3[-inTrain,]

set.seed(333)
inTrain <- createDataPartition(y=small4$classe, p=0.6, list=FALSE)
small_training4 <- small4[inTrain,]
small_testing4 <- small4[-inTrain,]
```

Random Forest (Test 1)

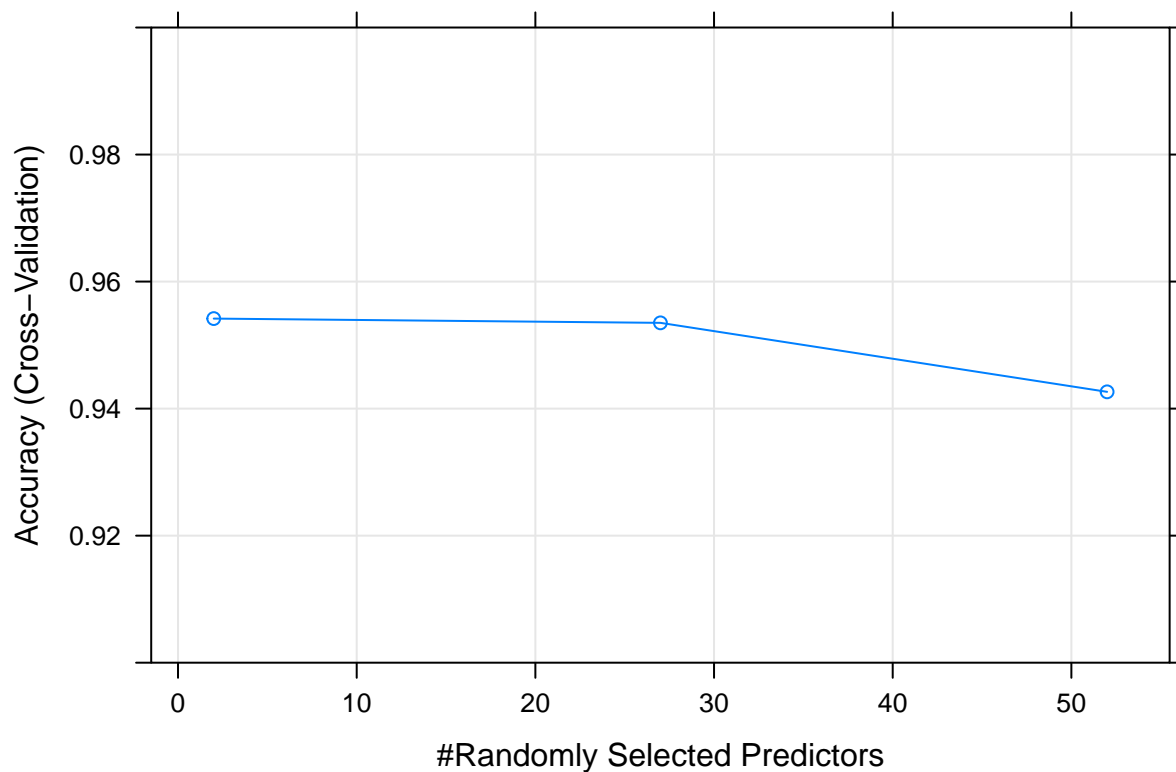
Selected random forest and run it on the first train data set using cross validation:

```
set.seed(2)
##Train
modFit <- train(small_training1$classe ~ ., method="rf", trControl=trainControl(method = "cv", number = 
print(modFit, digits=3)
```

```
## Random Forest
```

```
##
## 2946 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 2209, 2211, 2209, 2209
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.954 0.942 0.00527 0.00665
## 27 0.954 0.941 0.00554 0.00701
## 52 0.943 0.927 0.00782 0.00989
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

```
plot(modFit, ylim = c(0.9, 1))
```



```
##Test Set
predictions <- predict(modFit, newdata=small_testing1)
cf <- confusionMatrix(predictions, small_testing1$classe)
print(cf, digits=4)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 555  15    1    1    5
##           B   2 349  11    0    1
##           C   0  14 329  29    7
##           D   1   2   1 290    1
##           E   0   0   0   1 346
##
## Overall Statistics
##
##           Accuracy : 0.9531
##           95% CI : (0.9428, 0.962)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9406
##           McNemar's Test P-Value : 9.481e-08
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity       0.9946  0.9184  0.9620  0.9034  0.9611
## Specificity       0.9843  0.9911  0.9691  0.9970  0.9994
## Pos Pred Value    0.9619  0.9614  0.8681  0.9831  0.9971
## Neg Pred Value     0.9978  0.9806  0.9918  0.9814  0.9913
## Prevalence        0.2845  0.1938  0.1744  0.1637  0.1836
## Detection Rate    0.2830  0.1780  0.1678  0.1479  0.1764
## Detection Prevalence 0.2942  0.1851  0.1933  0.1504  0.1770
## Balanced Accuracy 0.9895  0.9548  0.9656  0.9502  0.9802
```

```
##Course Provided Test Set
print(predict(modFit, newdata=testing))
```

```
## [1] B A C A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

```
oos1 <- 1 - cf$overall[1]
```

Where overall accuracy results **0.9530852** and out of sample error is **0.0469148**

Random Forest (Test 2)

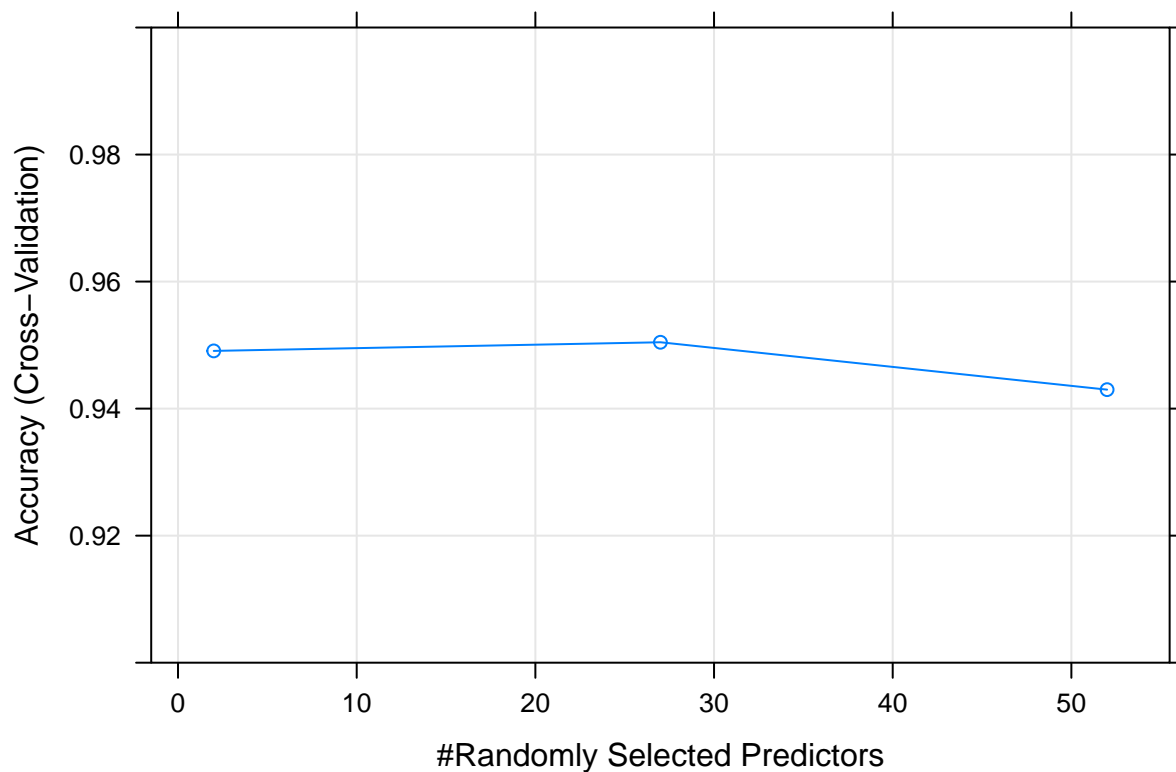
Second run adding reprocessing and cross validation:

```
set.seed(2)
##Train
modFit <- train(small_training1$classe ~ ., method="rf", preProcess=c("center", "scale"), trControl=trainControl(
print(modFit, digits=3)
```

```
## Random Forest
```

```
##
## 2946 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered (52), scaled (52)
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 2209, 2211, 2209, 2209
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.949 0.936 0.00600 0.00759
## 27 0.950 0.937 0.00701 0.00889
## 52 0.943 0.928 0.01046 0.01325
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.
```

```
plot(modFit, ylim = c(0.9, 1))
```



```
##Test Set
predictions <- predict(modFit, newdata=small_testing1)
cf <- confusionMatrix(predictions, small_testing1$classe)
print(cf, digits=4)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 555  11    0    1    0
##           B   3 354  15    0    4
##           C   0  13 323    9    3
##           D   0   1   3 309    1
##           E   0   1   1   2 352
##
## Overall Statistics
##
##           Accuracy : 0.9653
##           95% CI : (0.9562, 0.973)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9561
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9946  0.9316  0.9444  0.9626  0.9778
## Specificity      0.9914  0.9861  0.9846  0.9970  0.9975
## Pos Pred Value   0.9788  0.9415  0.9282  0.9841  0.9888
## Neg Pred Value   0.9978  0.9836  0.9882  0.9927  0.9950
## Prevalence       0.2845  0.1938  0.1744  0.1637  0.1836
## Detection Rate   0.2830  0.1805  0.1647  0.1576  0.1795
## Detection Prevalence 0.2891  0.1917  0.1775  0.1601  0.1815
## Balanced Accuracy 0.9930  0.9588  0.9645  0.9798  0.9876
```

```
##Course Provided Test Set
print(predict(modFit, newdata=testing))
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

```
oos2 <- 1 - cf$overall[1]
```

Where overall accuracy results **0.9653238** and out of sample error is **0.0346762**

Random Forest (Test 3)

Verify results using the second training dataset:

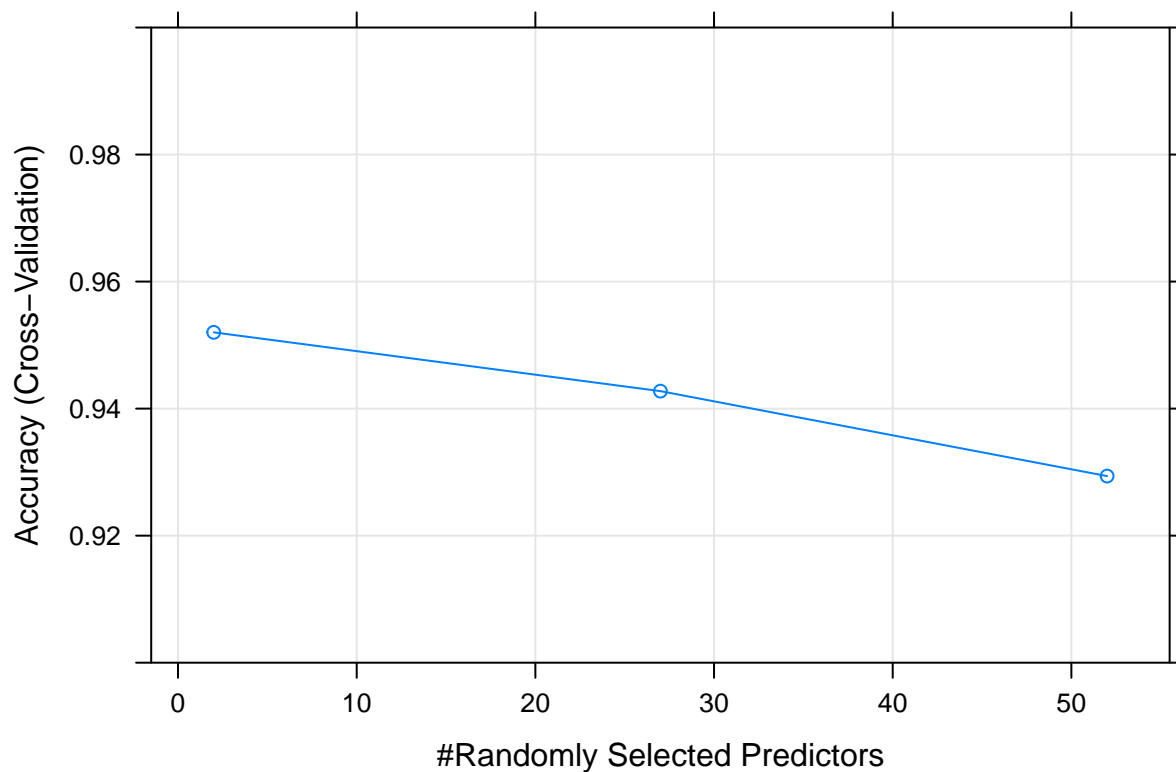
```
set.seed(2)
##Train
modFit <- train(small_training2$classe ~ ., method="rf", preProcess=c("center", "scale"), trControl=trainControl(
print(modFit, digits=3)
```

```
## Random Forest
```



```
##
## 2917 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered (52), scaled (52)
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 2188, 2188, 2188, 2187
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.952 0.939 0.00829 0.0105
## 27 0.943 0.928 0.01212 0.0153
## 52 0.929 0.911 0.01823 0.0231
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

```
plot(modFit, ylim = c(0.9, 1))
```



```
##Test Set
predictions <- predict(modFit, newdata=small_testing2)
cf <- confusionMatrix(predictions, small_testing2$classe)
print(cf, digits=4)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 548  26    0    1    0
##           B   1 337  18    2    1
##           C   0 11 316  20    6
##           D   3   0   4 294    5
##           E   0   2   0   1 345
##
## Overall Statistics
##
##           Accuracy : 0.948
##           95% CI : (0.9371, 0.9574)
##           No Information Rate : 0.2844
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9341
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9928  0.8963  0.9349  0.9245  0.9664
## Specificity      0.9806  0.9859  0.9769  0.9926  0.9981
## Pos Pred Value   0.9530  0.9387  0.8952  0.9608  0.9914
## Neg Pred Value   0.9971  0.9753  0.9861  0.9853  0.9925
## Prevalence       0.2844  0.1937  0.1741  0.1638  0.1839
## Detection Rate   0.2823  0.1736  0.1628  0.1515  0.1777
## Detection Prevalence 0.2962  0.1850  0.1819  0.1577  0.1793
## Balanced Accuracy 0.9867  0.9411  0.9559  0.9586  0.9822
```

```
##Course Provided Test Set
print(predict(modFit, newdata=testing))
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

```
oos3 <- 1 - cf$overall[1]
```

Where overall accuracy results **0.947965** and out of sample error is **0.052035**

Random Forest (Test 4)

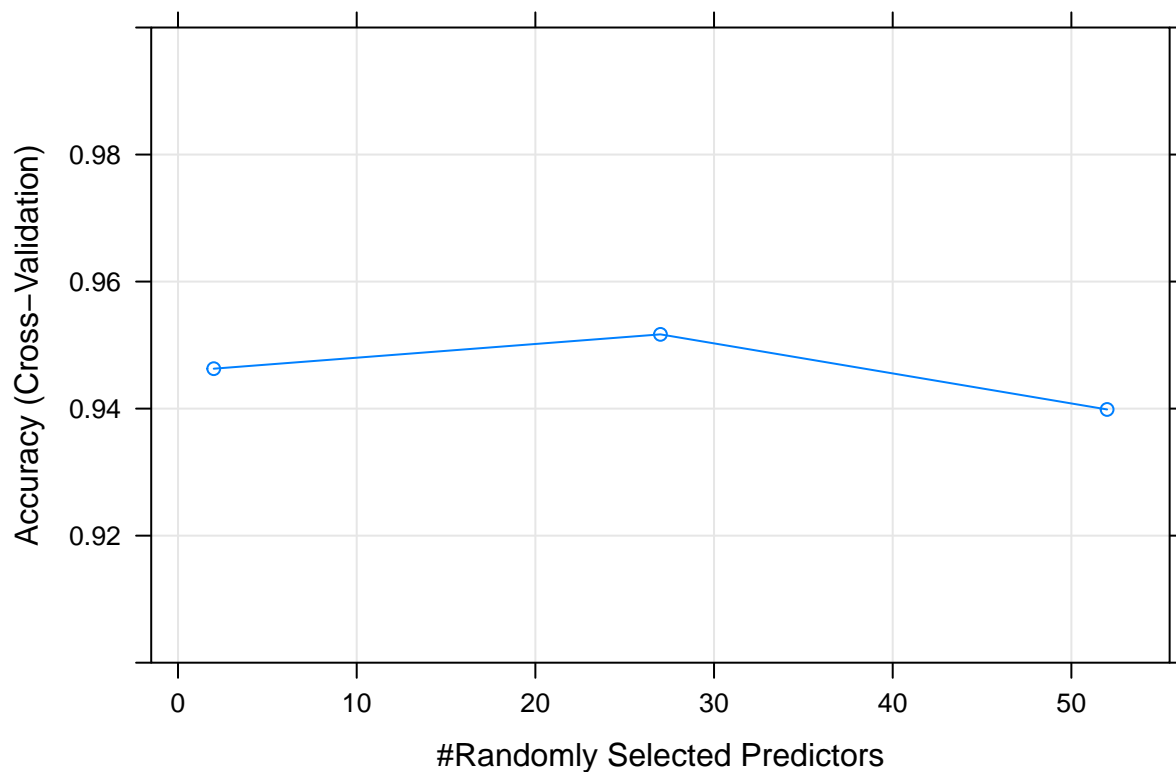
Last run on the third dataset as ulterior validation:

```
set.seed(2)
##Train
modFit <- train(small_training3$classe ~ ., method="rf", preProcess=c("center", "scale"), trControl=trainControl(
print(modFit, digits=3)
```

```
## Random Forest
```

```
##
## 2960 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered (52), scaled (52)
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 2220, 2220, 2220, 2220
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.946 0.932 0.0051 0.00642
## 27 0.952 0.939 0.0115 0.01453
## 52 0.940 0.924 0.0157 0.01987
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.
```

```
plot(modFit, ylim = c(0.9, 1))
```



```
##Test Set
predictions <- predict(modFit, newdata=small_testing3)
cf <- confusionMatrix(predictions, small_testing3$classe)
print(cf, digits=4)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 553  17    1    0    0
##           B   5 356    6    2    3
##           C   2   7 330    7    4
##           D   0   1   7 313    3
##           E   0   0   0   1 352
##
## Overall Statistics
##
##           Accuracy : 0.9665
##           95% CI : (0.9576, 0.974)
##           No Information Rate : 0.2843
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9576
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity       0.9875  0.9344  0.9593  0.9690  0.9724
## Specificity       0.9872  0.9899  0.9877  0.9933  0.9994
## Pos Pred Value    0.9685  0.9570  0.9429  0.9660  0.9972
## Neg Pred Value     0.9950  0.9844  0.9914  0.9939  0.9938
## Prevalence        0.2843  0.1934  0.1746  0.1640  0.1838
## Detection Rate     0.2807  0.1807  0.1675  0.1589  0.1787
## Detection Prevalence 0.2898  0.1888  0.1777  0.1645  0.1792
## Balanced Accuracy  0.9874  0.9622  0.9735  0.9812  0.9859
```

```
##Course Provided Test Set
print(predict(modFit, newdata=testing))
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

```
oos4 <- 1 - cf$overall[1]
```

Where overall accuracy results **0.9664975** and out of sample error is **0.0335025**

Random Forest (Test 5)

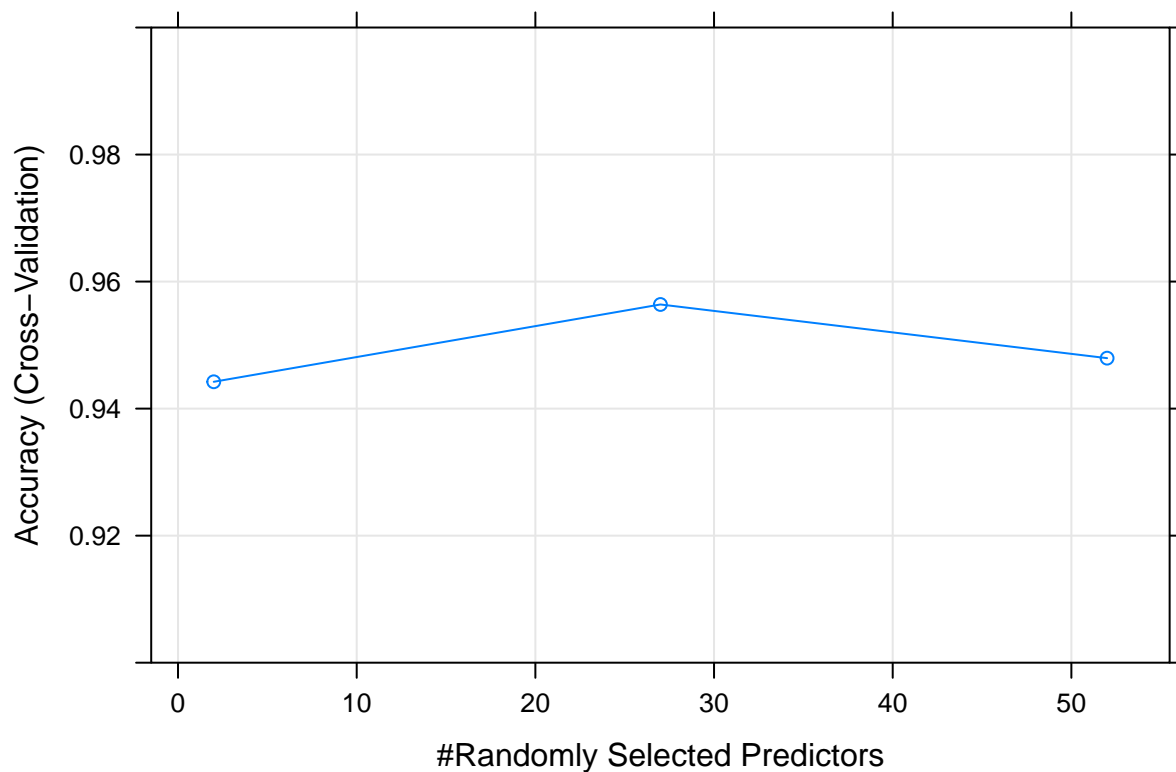
Final run on fourth dataset using preprocess and cross validation:

```
set.seed(2)
##Train
modFit <- train(small_training4$classe ~ ., method="rf", preProcess=c("center", "scale"), trControl=trainControl(
print(modFit, digits=3)
```

```
## Random Forest
```

```
##
## 2958 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## Pre-processing: centered (52), scaled (52)
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 2217, 2219, 2219, 2219
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.944 0.929 0.00712 0.00907
## 27 0.956 0.945 0.00756 0.00959
## 52 0.948 0.934 0.00810 0.01029
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.
```

```
plot(modFit, ylim = c(0.9, 1))
```



```
##Test Set
predictions <- predict(modFit, newdata=small_testing4)
cf <- confusionMatrix(predictions, small_testing4$classe)
print(cf, digits=4)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 549  10    0    0    0
##           B   5 357  11    0    1
##           C   1 12 325    8    0
##           D   5   0   7 315   10
##           E   0   2   0   0 351
##
## Overall Statistics
##
##           Accuracy : 0.9634
##           95% CI : (0.9542, 0.9713)
##           No Information Rate : 0.2844
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9538
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9804  0.9370  0.9475  0.9752  0.9696
## Specificity      0.9929  0.9893  0.9871  0.9866  0.9988
## Pos Pred Value   0.9821  0.9545  0.9393  0.9347  0.9943
## Neg Pred Value   0.9922  0.9850  0.9889  0.9951  0.9932
## Prevalence       0.2844  0.1935  0.1742  0.1640  0.1838
## Detection Rate   0.2788  0.1813  0.1651  0.1600  0.1783
## Detection Prevalence 0.2839  0.1899  0.1757  0.1712  0.1793
## Balanced Accuracy 0.9866  0.9632  0.9673  0.9809  0.9842
```

```
## Course Provided Test Set
print(predict(modFit, newdata=testing))
```

```
## [1] B A B A A E D B A A B C B A E E A B A B
## Levels: A B C D E
```

```
oos5 <- 1 - cf$overall[1]
```

Where overall accuracy results **0.9634332** and out of sample error is **0.0365668**

Out of Sample Error Rate

The average of the sample error rates generated by the random forest method using pre-processing and cross validation on the 5 test sets provided a predicted out of sample rate of **0.0407391**.